

Question 1. Case Study: Making a Brand Discoverable in the Age of LLMs Marks : 10.00

Live Proctored Assessment - Brand Content & Strategy Manager

Format: This is a live, proctored exercise, completed in a single supervised session. Work independently — no outside resources, collaborators, or AI assistants (ChatGPT, Gemini, Claude, Copilot, etc.). We are assessing how you reason.

Response: Bullet points are welcome — we're grading your reasoning, not your prose, so don't worry about polished English or framing full sentences. Aim for a focused answer across Q1–Q4 combined (roughly 400–800 words total, or equivalent bullets); a sharp short answer beats a long rambling one.

1. Problem Statement

You have joined a consumer brand as its Brand Content & Strategy Manager. For two decades, "helping people find us" meant Search Engine Optimization. That world is shifting: people increasingly ask an assistant (ChatGPT, Gemini, Claude, Perplexity, Copilot) instead of searching. The assistant replies in one paragraph, names a few brands or none, and the user rarely clicks a link. If your brand isn't named in that paragraph, for that user it doesn't exist. And even when a user does click through, many bounce off the site within seconds — unable to quickly find what they came for, they leave without ever really engaging. This is now a technical problem wearing a marketing costume — which is why we're handing it to engineers.

unable to quickly find what they came for, they leave without ever really engaging. This is now a technical problem wearing a marketing costume — which is why we're handing it to engineers.

2. The symptoms we're hitting

- **Invisible in AI answers.** Ask an AI app such as ChatGPT or Gemini about our category and we're not named at all — it lists competitors and omits us entirely.
- **Stale facts.** Where we do appear, the details are sometimes two years old (old logo, discontinued product, wrong price). The fresh facts don't seem to be what the assistant finds and repeats.
- **Visitors don't stay engaged.** Even when users do land on our site, they bounce quickly. They arrive without the context an assistant's summary gave them, feel lost navigating the site, and can't easily find the specific information they were looking for — so they leave. The site also treats every visitor identically: it doesn't retain or adapt to the context and intent a person arrived with, so there's no continuity from the assistant's answer to the page and no personalization to keep them engaged.

A and B share a common underlying reason: machines are now the first readers of our content, and our content was written for humans, not machines. C stems from a related but distinct gap: our site itself isn't designed to orient a visitor who arrives mid-journey, already primed with a specific question.

(The Appendix at the end explains the general concepts you'll need — how discovery and AI citation actually work — with no reference to this scenario. Read it before answering.)

3. The example: Garuda Footwear (fictional)

Picture a large, fictional Indian shoe brand, Garuda Footwear — founded 1994, 400+ stores, flagship Garuda Velocity running shoes (new model: Velocity X9), plus cricket, kids', and ethnic lines. It was rebranded in 2024 (new logo, tagline "Built to Fly", recycled soles).

Its failures, as leadership sees them:

- **(1) Invisible.** "best Indian running shoes under 6000" on AI apps such as ChatGPT and Gemini lists rivals and omits Garuda entirely — it simply isn't named.
- **(2) Stale facts.** Where mentioned, an assistant repeats the old 2019 logo, the dead tagline, and a discontinued shoe's price — the new facts aren't repeated widely enough across the web, so what gets found and cited is still the old version.
- **(3) Visitors bounce.** A user asks an assistant about the Velocity X9, clicks through to Garuda's site, and lands on the generic homepage — no trace of the specific question that brought them there. They can't quickly find the X9's page, price, or availability, feel lost, and leave within seconds. The site shows the same generic page to everyone: it doesn't recognize or retain the visitor's context, so a user primed with a specific question (or a returning one) gets no personalized continuity.

4. Scope — what we want, and don't (read before answering)

We want **TECHNICAL and CREATIVE ideas**: crawlers, indexing, structured data/markup, site rendering, retrieval-friendliness, machine-readable content, entity identity, and monitoring; plus creative tactics for making facts consistent, corroborated, and quotable across the web.

We do **NOT** want business/commercial ideas. No paid partnerships, sponsorships, or influencer deals; no paying an LLM vendor for placement/ads; no agency retainers, hiring, M&A, or co-branding. If your main lever is "spend money so someone else features us," it is out of scope. Reason like an engineer who read the manual, not a sales team booking ad slots.

5. Your task

Look at the failures listed above. Answer the questions below to help Garuda's leadership. Each of Q1–Q3 has two parts — (a) identify the root cause and (b) suggest a fix. Bullet points are fine.

Q1 — Visibility (invisible in AI apps such as ChatGPT and Gemini)

- **(a) Root cause.** Explain the underlying reasons why Garuda is absent from AI apps such as ChatGPT and Gemini. Reason from how these systems actually work — don't just re-narrate the symptom.
- **(b) Fix.** Propose concrete technical and creative steps that would make Garuda discoverable and named by these assistants.

Q2 — Staleness (outdated / wrong facts)

- **(a) Root cause.** Explain why the few facts that do surface (e.g. on a direct-name query) are stale — old 2019 logo, dead tagline, discontinued shoe's price. Reason about what gets found and repeated across the web, and why the fresh 2024 facts don't.
- **(b) Fix.** Propose concrete technical and creative steps that would make the current facts the ones assistants find, trust, and repeat.

Q3 — Engagement (visitors bounce; no context retention / personalization)

- **(a) Root cause.** Explain why visitors who do land on the site bounce off without engaging. Reason about what happens (or doesn't) between an assistant's answer and the on-site experience — including whether the site retains and adapts to the context and intent a visitor arrived with, or treats every visitor identically.
- **(b) Fix.** Propose concrete technical and creative steps that would help visitors orient quickly and find what they came for, reducing bounce — including how the site could retain and use a visitor's context (e.g. the question or intent they arrived with, or prior visits) to personalize the experience and keep them engaged.

Q4 — Spot a symptom we didn't see.

The three symptoms above (Invisible / Stale Facts / Visitors Bounce) are not exhaustive. Identify another distinct way Garuda fails in an assistant-first world — one that isn't a restatement of visibility, staleness, or bounce — name it in one line, explain its root cause from how these systems work, propose a concrete technical/creative fix, and note how you'd measure it. One symptom, done well, beats several shallow ones. Same scope rules apply (no business/ad levers).

For each of Q1–Q3, briefly note how Garuda would know the fix is working — how you'd monitor visibility, factual accuracy, and on-site engagement across multiple assistants and visits over time.

A focused answer with a few excellent, well-targeted ideas beats a shallow list of fifteen. State any assumptions in one line. No code or tools required.

6. Rubric

Criterion	Looking for
Root-cause understanding (part a of Q1–Q3)	Correctly explains why the symptoms happen; reasons about retrieval, crawlers, page structure, and corroboration
Technical measures (part b of Q1–Q3)	Specific, plausible, correctly targeted; names real mechanisms and the causal link to the root cause

Creative problem identification & solution (Q4)	Identifies a genuinely distinct, in-scope symptom beyond Invisible / Stale / Bounce, with a sound root cause and a non-obvious, correctly targeted fix
Creative problem identification & solution (Q4)	Invisible / Stale / Bounce, with a sound root cause and a non-obvious, correctly targeted fix
Measurement & rigor	A sane way to test/monitor visibility & accuracy across assistants over time
Use of example & clarity	Grounded in Garuda; clear enough for a non-technical leader
Scope discipline	Stays technical/creative; avoids business/ad-buying levers

We care less about whether you know the buzzwords, and more about how clearly you reason from how these systems work to what a brand can concretely do about it.

Appendix — Background concepts

This appendix explains the general ideas behind the exercise, with no reference to any specific company. It describes how these systems behave — not what to do about any particular problem. Working out the fixes is your job. A glossary of key terms appears at the end of this document.

A. How search visibility works (the basics).

Search engines, and increasingly AI assistants, discover content the same basic way: an automated program (a "crawler") visits a page, reads what's there, and stores it so it can be found later. For a page to be visible at all, three things have to succeed in order — the crawler has to be let in, it has to be able to read what's on the page, and it has to be able to pick out the specific fact someone is looking for. If any one of those steps fails, the page effectively doesn't exist for that system — plainly visible to a human, yet invisible to the machine.

B. How assistants like ChatGPT use sources.

When you ask an assistant a question, it doesn't always answer purely from memory. Many assistants can go and look things up in the moment — searching, fetching a few pages, and building the answer from what those pages happen to say, sometimes showing links back to them. What ends up in the answer depends on what the assistant can find and use at that

moment: the pages that get picked as sources tend to be the ones a machine could easily reach, easily read, and easily quote a clear fact from. If nothing about a brand is easy to reach, read, and quote, it simply won't make it into the answer at all.

C. How machines "read" a page.

A page that looks complete to a person isn't always complete to a machine. Some content is only assembled after a page loads, or lives in a form a simple reader can't interpret — so a fact that's plainly visible on screen can be invisible to a program that isn't looking at it the same way. In general, the more explicitly and unambiguously a fact is stated in plain, readable text, the more likely a machine is to extract it correctly; the more it's implied, buried, or locked inside something non-textual, the more likely it's missed.

D. Why agreement across the web matters.

Machines tend to treat a fact as more trustworthy when many independent places say the same thing. A claim that lives in only one spot is fragile; a claim repeated consistently across lots of unrelated sources is far more likely to be believed and repeated back. A related problem is mistaken identity: when several different things share a name, a system can mix them up unless there's something that clearly distinguishes one from the others. So how a brand is described across the wider web — not just on its own pages — shapes what assistants say about it.

E. Personalization and prior context.

These assistants don't treat every user identically. They can draw on context they already hold about a person — earlier messages in the conversation, past questions, stated preferences, location, or other signals — and use it to tailor what they surface, so the answer is weighted toward what seems most relevant to that user. This is one more reason two people can ask the very same question and get different brands, framing, or emphasis in return: part of the answer depends not just on the content that exists out there, but on who's asking.

H. Discovery is a funnel, not an on/off switch.

A page can pass one stage and fail at the next:

Exists→discovered→accessible→readable→understood→retrieved for this question→selected as a source→reflected in the answer

So "the page is indexed" does not necessarily mean "the assistant will mention it." Different failures can happen at different stages.

J. The web can contain several versions of the same fact.

Updating a company's own page does not erase older versions elsewhere. An assistant might encounter an old article, retailer page, cached description, directory listing, newer official page, and third-party review that all say slightly different things. It then has to decide which evidence to rely on. Recency, clarity, source quality, consistency, and agreement between sources can all affect which version survives into the answer.

Glossary

Term	Plain-English meaning
Crawler	Software that automatically visits and reads webpages.
Index	A stored representation of content that a system can later find.
Retrieval	Finding information relevant to a particular question.
Rendering	Turning a webpage's code into the content a visitor — or machine — can actually read.
Structured data	Machine-readable labels describing what things on a page mean.
Entity	A distinct real-world thing: a company, product, person, place, etc.
Citation / source	A page or document an assistant relies on when constructing an answer.

Case Study Solution: Making a Brand Discoverable in the Age of LLMs

Role: Brand Content & Strategy Manager

Company: Garuda Footwear (Fictional)

Executive Summary

The transition from traditional Search Engine Optimization (SEO) to AI Assistant Optimization requires addressing machine readability, web-wide factual consensus, and mid-journey landing continuity. This strategy outlines technical and creative measures to make Garuda Footwear discoverable, accurately cited, and engaging for AI-driven shoppers without relying on paid advertising or sponsored partnerships.

Q1 — Visibility: Addressing Absence in AI Answers

(a) Root Cause

AI systems rely on real-time retrieval crawlers (e.g., GPTBot, PerplexityBot) to construct answers. Garuda suffers from a rendering and entity-linkage failure at the retrieval stage: its product catalog relies heavily on dynamic JavaScript without semantic microformats. Crawlers cannot extract plain-text facts directly connecting the entity **"Garuda Footwear"** to **"Indian running shoes under ₹6000,"** causing LLMs to select competitor pages that are easily parsed.

(b) Proposed Fix

- **Technical:** Implement Server-Side Rendering (SSR) across all catalog and category pages to expose clean, static HTML. Include JSON-LD schemas (Brand, Product, Offer) to explicitly define brand category and price bounds in machine-readable labels.
- **Creative:** Publish an open-access, text-rich web directory titled *"Official Technical Buyer's Guide for Indian Running Shoes Under ₹6000"* featuring simple, unnested HTML comparison tables that retrieval crawlers can easily extract and quote.

Q2 — Staleness: Overcoming Legacy Brand Data

(a) Root Cause

LLMs judge factual truth via cross-web consensus and corroboration. Legacy 2019 data (old logo, dead tagline, discontinued shoe prices) exists across hundreds of third-party review sites, blogs, and cached directories. Because the 2024 rebrand facts live almost exclusively

on Garuda's primary domain, search algorithms weight the heavily repeated legacy data as the prevailing ground truth.

(b) Proposed Fix

- **Technical:** Embed `dateModified`, `datePublished`, and `sameAs` entity markup across all primary web pages. Set HTTP 301 permanent redirects for deprecated product URLs and submit updated XML sitemaps to major search and retrieval indexes.
- **Creative:** Distribute an open digital media kit and standardized structured data feeds to online footwear directories, encyclopedias, and consumer product databases to establish machine-verifiable consensus across independent web sources.

Q3 — Engagement: Reducing Visitor Bounce Rates

(a) Root Cause

High bounce rates stem from a contextual gap between the assistant's response and the website's entry experience. The site operates statelessly, serving a generic homepage regardless of how or why the user arrived. Visitors primed with a specific query by an AI assistant (e.g., *"Velocity X9 price and specs"*) land mid-journey on a generic page without context retention, feel lost navigating, and leave within seconds.

(b) Proposed Fix

- **Technical:** Parse referral HTTP headers and incoming URL query parameters (e.g., `?utm_intent=velocity_x9`) to dynamically route traffic or trigger contextual UI overlays, storing session intent via `localStorage`.
- **Creative:** Implement a top-of-page contextual entry banner (*"Looking for the Velocity X9?"*) on arrival pages that immediately surfaces key product specifications, pricing, and live size availability matching the visitor's entry intent.

Q4 — Distinct Symptom: Category & Attribute Misclassification

(a) Identified Symptom

Product Usage & Category Misclassification: AI assistants incorrectly label the flagship Velocity X9 as a "trail running shoe for mud" instead of a road trainer, generating improper usage, buyer dissatisfaction, and negative product reviews.

(b) Root Cause

A lack of explicit negative and terrain metadata on product pages forces LLM probabilistic models to guess shoe usage characteristics based on loose, descriptive marketing copy.

(c) Proposed Fix

Include explicit, plain-text key-value attribute tables in HTML schema (terrain: Road, waterproof: False) directly within product page structures.

(d) Measurement Metric

Run automated weekly prompt scripts across ChatGPT, Gemini, Claude, and Perplexity querying feature specifics against a master product specifications matrix to calculate a Category Accuracy Percentage.

Performance & Monitoring Strategy (Q1–Q3)

- **Visibility & Factual Accuracy:** Execute automated weekly query scripts across major AI platforms targeting high-intent category keywords to monitor brand mention frequency and verify that cited facts (logo, tagline, current pricing) align with current data.
- **On-Site Engagement:** Track AI referral traffic channels (chatgpt.com, perplexity.ai, claude.ai) in web analytics to measure average session depth, bounce rate reductions, and checkout conversions.